


Chapter 4

Processes and Threads

- **Resource Ownership**
 - Includes a virtual address space that hold the process image(PCB)
 - OS performs a protection function to prevent interference between processes resources
- **Scheduling/Execution**
 - Execution path that maybe interleaved with other processes
 - Process has an execution state and a dispatching priority and that is controlled by the OS
- **Thread or Lightweight Process:** Unit of dispatching
- **Process or Task:** The unit of resource ownership
- **Multi-threading:** Ability of the OS to support multiple, concurrent paths of execution within a single process


Single Threaded Approaches

- Single thread of execution per process is referred to as a single threaded approach
- MS-DOS is an example



Multi-threaded Approach

- A Java run-time environment is an example of a system of one process with multiple threads



Process

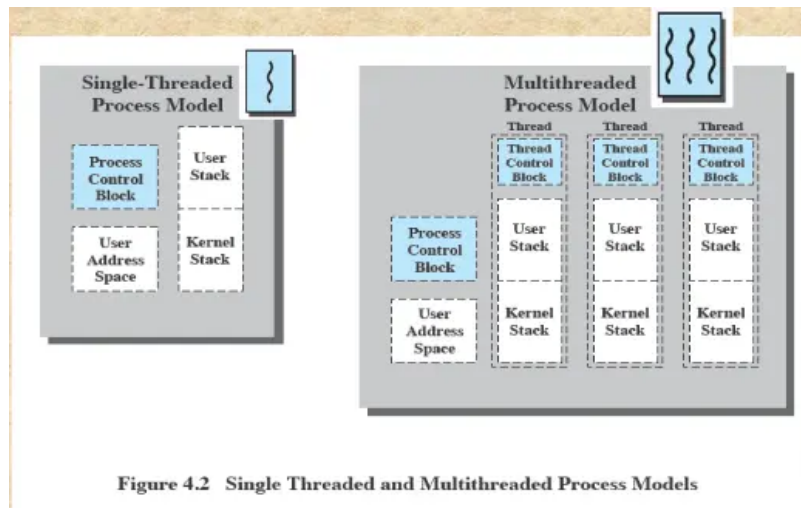
- The unit of *resource allocation* and a unit of *protection*
- A virtual address space that holds the process image
- Protected access to:
 - Processors

- Other processes
- files
- I/O resources

One of More Threads in a Process

- Each Thread has:
 - execution state
 - saved thread context when not running
 - execution stack
 - per-thread static storage for local variables
 - access to the memory and resources of its process (shared amongst threads)

Threads vs Processes



Benefits of Threads

- Takes less time to create a new thread
- Less time to terminate a thread
- Switching between 2 threads takes less time
- Threads enhance efficiency in communication between programs

Thread use in a single-user system

- Foreground and Background work
- Asynchronous processing
- Speed of execution
- Modular program structure

Threads


- Scheduling and dispatching is done on a thread basis on the OS that supports threads
- State information dealing with execution is stored in a thread level data structure
 - *suspending* a process involves suspending all threads of the process
 - *terminating* of a process terminates all threads within the process

Thread Execution State

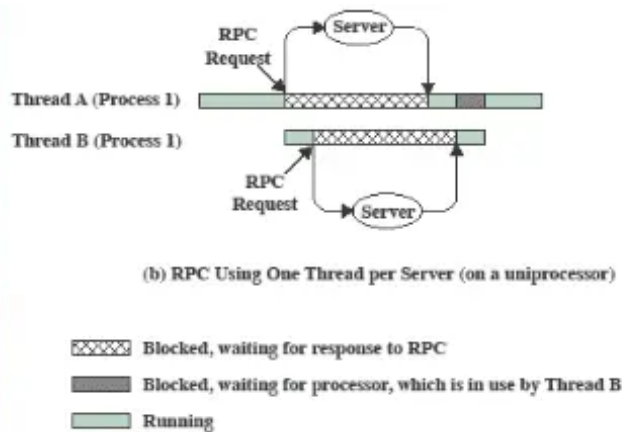
- Key states
 - *Running*
 - *Ready*
 - *Blocked*

- Thread operations associated with a change in thread state are
 - *Spawn*
 - *Block*
 - *Unblock*
 - *Finish*

RPC Using Single Thread



RPC Using One Thread per Server



Multi-threading on a Uni-processor




Figure 4.4 Multithreading Example on a Uniprocessor

Thread Synchronization

- Necessary to sync activities of the various threads
 - all threads of a process share the address space and other resources
 - any alteration will affect the other threads in the same process

Types of Threads

• User Level Thread

- All thread management is done by the application
- Kernel not aware of its existence

- Advantages

- Thread switching does not require kernel mode privileges

- Scheduling can be application specific

- ULT's can run on any OS

- Disadvantages*

- System calls are blocked which in turn also blocks all other threads within the process

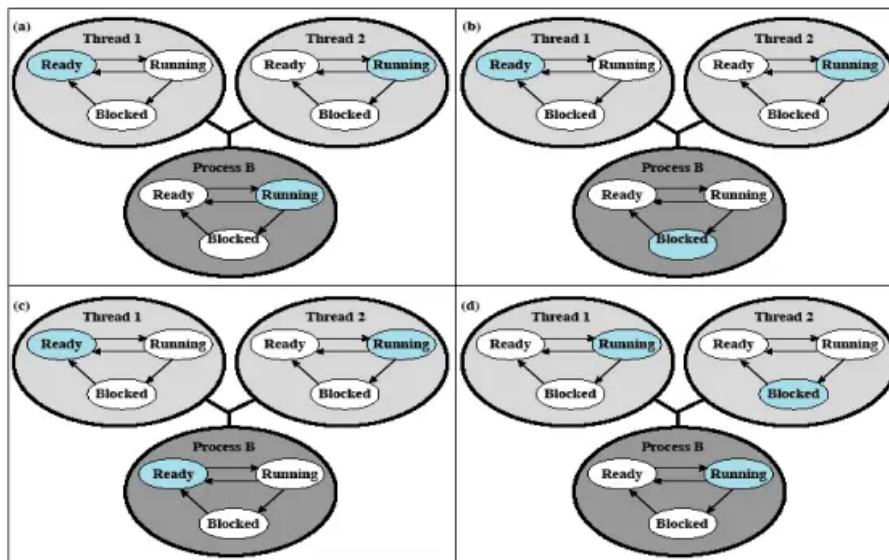
- A multi-threading application cannot take advantage of multiprocessing if it is purely ULT

- Overcoming ULT Disadvantages

- converts a blocking system call into a non-blocking system call - jacketing

- writing an application as multiple processes rather than multiple threads

- Relation between ULT states and Process States*



• Kernel Level Thread

- Thread management is done by the kernel
 - no thread management is done by the application i.e. windows

• Advantages of KTLs

- Kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread is blocked, the kernel reschedules another thread of the same process
- routines can be multithreaded

• Disadvantage

- Transfer of control from one thread to another within the same process requires mode switch to the kernel

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Combined Approaches (i.e. Solaris)

- Thread creation is done in the user space
- Bulk of scheduling and syncing is done by application

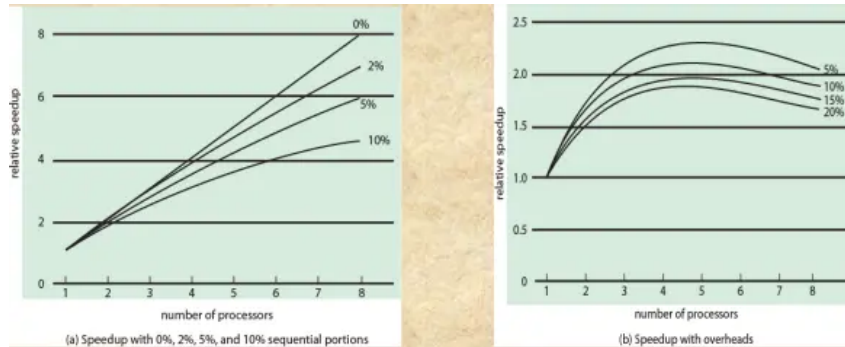
Relation between threads and processes

- 1:1 : Each thread of execution is a unique process with its own address space and resources. Unix


- *M:1*: Defined address space and dynamic resource ownership. Multiple threads created and executed. Windows, Linux
- *1:M*: Migrated thread from one process to another. Allows a thread to be easily moved among different systems. Ra, Emerald
- *M:N*: Combination of *M:1* and *1:M*. TRIX

Multi-threading and Multicore

$$\text{Speedup} = \frac{\text{time to execute on 1 processor}}{\text{time to execute on } N \text{ parallel processors}} = \frac{1}{(1-f) + \frac{f}{N}}$$



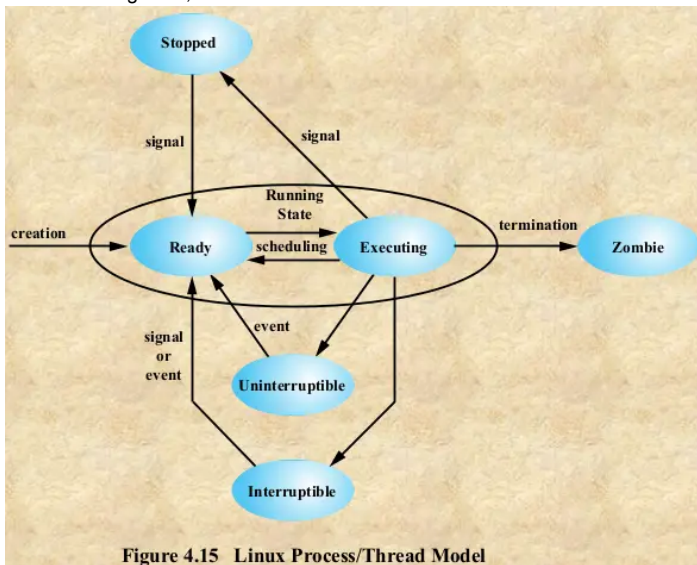
Database Workloads on Multiple-Processor Hardware



Linux Tasks

- A process or task in Linux is represented by a `task_struct` data structure. It contains
 - *State* => executing, ready, blocked
 - *Scheduling Info* => priority, time slice allowed
 - *Identifiers* => PID, userID, groupID
 - *IPC*
 - *Links* => to parent, sibling, children
 - *Times/Timers* => processor time used so far, interval timer
 - *File system* => pointers to opened files, current directory of process
 - *Address Space* => program, data

- *Context* => registers, stack



Linux Threads

- Linux does not recognise a distinction between threads and processes
- New process is created by copying the attributes of the current process (fork or clone)
- new process can be cloned so that it shares resources (address space (VM), signal handlers, files, IO etc)
- Processes sharing same VM operate as threads within a single process
- Both clone & pthreads make what text calls KLTs